

Modules

Module – 1 Introduction to MongoDB

- 1.1 Overview of MongoDB**
- 1.2 NoSQL vs. SQL databases**
- 1.3 Advantages of using MongoDB**
- 1.4 Installing and setting up MongoDB**

Module – 2 MongoDB Data Model

- 2.1 Understanding Documents and Collections**
- 2.2 BSON Data Format**
- 2.3 Document-Oriented Data Model**
- 2.4 MongoDB Schema Design Best Practices**

Module – 3 CRUD Operations in MongoDB

- 3.1 Creating documents**
- 3.2 Reading documents**
- 3.3 Updating documents**
- 3.4 Deleting documents**
- 3.5 Querying data with MongoDB**

Module – 4 Indexing and Query Optimization

- 4.1 Importance of indexing**
- 4.2 Creating and Managing Indexes**
- 4.3 Query Optimization Techniques**
- 4.4 Using the Aggregation Framework**

Module – 5 Data Modeling in MongoDB

- 5.1 Embedded vs. Referenced documents**
- 5.2 Data normalization and de-normalization**
- 5.3 Modeling Relationship**

Module – 6 Advanced MongoDB Features

- 6.1 Geospatial Queries**
- 6.2 Text search**
- 6.3 Full-Text Search with Text Indexes**
- 6.4 Time-Series Data with MongoDB**

Module – 7 MongoDB Atlas and Cloud Deployment

- 7.1 Introduction to MongoDB Atlas**
- 7.2 Creating and Managing Clusters**
- 7.3 Deploying MongoDB in the Cloud**

Module – 8 Security and Authentication

- 8.1 Securing MongoDB**
- 8.2 User Authentication and Authorization**
- 8.3 Role-Based Access Control (RBAC)**

Module – 9 Backup and Recovery

- 9.1 Backup Strategies**
- 9.2 Restoring Data**
- 9.3 Disaster Recovery Planning**

Module – 10 MongoDB Aggregation Pipeline

- 10.1 Aggregation Concepts**
- 10.2 Using Pipeline Stages**
- 10.3 Custom Aggregation Expressions**

Module – 11 MongoDB Drivers and APIs

- 11.1 MongoDB Drivers**
- 11.2 Programming Language of your Choice**

Module – 12 MongoDB Performance Tuning

- 12.1 Profiling and Monitoring**
- 12.2 Query Performance Optimization**
- 12.3 Hardware and Resource Considerations**

Module – 13 Replication and Sharding

- 13.1 Replication for High Availability**
- 13.2 Sharding for Horizontal Scaling**
- 13.3 Configuring Replica Sets and Sharded Clusters**

Module – 14 Working with GridFS

- 14.1 Storing Large files in MongoDB**
- 14.2 Using GridFS for file Management**

4

Indexing and Query Optimization

4.1 Importance of Indexing

Indexes play a pivotal role in database systems, including MongoDB, as they significantly enhance the speed of query execution. Without indexes, MongoDB would need to perform a collection scan, which involves scanning every document in a collection to locate matching documents. This can be slow and resource-intensive, especially for large collections. Indexes work by providing a fast and efficient way to look up data based on specific fields, significantly reducing query response times.

Benefits of indexing in MongoDB include:

- 1. Faster Query Performance:** Indexes allow MongoDB to quickly locate and retrieve documents that match query criteria, resulting in reduced query execution times.
- 2. Reduced Resource Usage:** Indexed queries use fewer system resources (CPU, memory, disk I/O) compared to full collection scans, making your application more efficient.

3. Enforcement of Uniqueness: Unique indexes can enforce data integrity by ensuring that specific fields have unique values, preventing duplicate entries.

4. Support for Sorting: Indexes can be used for sorting query results, improving the efficiency of queries that involve sorting.

5. Covered Queries: Indexes can make certain queries "covered," meaning all required fields are in the index itself, eliminating the need to access the actual documents.

4.2 Creating and Managing Indexes

MongoDB provides various options for creating and managing indexes on your collections.

Creating Single Field Index

To create an index on a single field, you can use the "createIndex()" method:

- javascript

// Create an index on the "name" field of the "products" collection

```
db.products.createIndex({ name: 1 });
```

Here, “{ name: 1 }” specifies an ascending index on the "name" field. You can use “-1” for descending indexes.

Creating Compound Index

Compound indexes involve multiple fields and can significantly improve query performance for queries that filter or sort based on these fields. For example:

- javascript

// Create a compound index on the "category" and "price" fields

```
db.products.createIndex({ category: 1, price: 1 });
```

Creating Unique Index

Unique indexes enforce uniqueness constraints on specific fields. Attempts to insert duplicate values in the indexed field will result in an error:

- javascript

// Create a unique index on the "email" field of the "users" collection

```
db.users.createIndex({ email: 1 }, { unique: true });
```

Managing Indexes

You can view existing indexes, drop indexes, or rebuild them using MongoDB's index management methods:

- **javascript**

// List all indexes for a collection

```
db.products.getIndexes();
```

// Drop an index

```
db.products.dropIndex("name_1");
```

// Rebuild all indexes for a collection

```
db.products.reIndex();
```

It's important to design indexes that align with your application's query patterns. Over-indexing can lead to increased storage requirements and slower write operations, so strike a balance between query performance and storage overhead.

4.3 Query Optimization Techniques

MongoDB offers several techniques for optimizing queries:

Explain Method

The “explain()” method helps you understand how MongoDB executes a query. It provides information about the query execution plan, including which indexes are used, the number of documents examined, and the execution time.

- javascript

// Explain the execution plan for a query

```
db.products.find({ category: "Electronics"
}).explain("executionStats");
```

Reviewing the output of “explain()” can help identify areas for query optimization.

Covered Queries

A query is considered "covered" when all the fields needed to satisfy the query are included in the index itself. Covered queries can be significantly faster because MongoDB doesn't need to access the actual documents.

- javascript

// Create an index that covers the query

```
db.products.createIndex({ category: 1, price: 1 });
```


// Perform a covered query

```
db.products.find({ category: "Electronics" }, { _id: 0, name: 1, price: 1 });
```

Limit and Skip

Use the `limit()` and `skip()` methods to control the number of documents returned by a query. Be cautious with `skip()` as it can be inefficient for large result sets.

- javascript

// Limit the number of results

```
db.products.find().limit(10);
```

// Skip the first 5 results and then limit to 10

```
db.products.find().skip(5).limit(10);
```

Index Hinting

You can use the `hint()` method to explicitly specify which index to use for a query. This can be useful when you want to ensure a specific index is utilized.

- javascript

// Use the "category" index for the query

```
db.products.find({ category: "Electronics" }).hint({ category: 1 });
```

4.4 Using the Aggregation Framework

MongoDB's Aggregation Framework is a powerful tool for performing complex data transformations and aggregations. It allows you to filter, group, project, and compute data across documents in a collection. The Aggregation Framework is particularly useful when you need to perform operations like summarization, joining, and statistical analysis.

Here's an example that demonstrates the Aggregation Framework to calculate the average price of products in each category:

- javascript

// Calculate the average price of products in each category

```
db.products.aggregate([
  {
    $group: {
      _id: "$category",
      avgPrice: { $avg: "$price" }
    }
  },
])
```

```
{
  $project: {
    category: "$_id",
    _id: 0,
    avgPrice: 1
  }
}
]);
```

In this example:

- The “\$group” stage groups products by category and calculates the average price within each group.
- The “\$project” stage reshapes the result to include only the “category” and “avgPrice” fields.

The Aggregation Framework provides a rich set of operators and stages for data manipulation, making it a valuable tool for complex data processing tasks.

Example:

Let's consider a scenario where you have a collection called "orders" that stores information about customer orders. You want to find the total value of orders placed by each customer. Using the Aggregation Framework, you can achieve this:

- javascript

// Sample "orders" collection

```
[
  {
    "_id": ObjectId("5fd453fb4e0c103f9cb7f981"),
    "customer_id": ObjectId("5fd453fb4e0c103f9cb7f97d"),
    "order_date": ISODate("2022-01-15T09:30:00Z"),
    "total_amount": 200.0
  },
  {
    "_id": ObjectId("5fd453fb4e0c103f
9cb7f982"),
    "customer_id": ObjectId("5fd453fb4e0c103f9cb7f97e"),
    "order_date": ISODate("2022-02-20T14:15:00Z"),
    "total_amount": 150.0
  },
  {
    "_id": ObjectId("5fd453fb4e0c103f9cb7f983"),
    "customer_id": ObjectId("5fd453fb4e0c103f9cb7f97d"),
    "order_date": ISODate("2022-03-10T17:45:00Z"),
```

```
"total_amount": 300.0
}
]
```

To find the total value of orders placed by each customer, you can use the Aggregation Framework:

- javascript

```
db.orders.aggregate([
{
  $group: {
    _id: "$customer_id",
    totalOrderValue: { $sum: "$total_amount" }
  }
},
{
  $lookup: {
    from: "customers",
    localField: "_id",
    foreignField: "_id",
    as: "customer_info"
  }
}
```

```
},  
  
{  
  $project: {  
    customer_id: "$_id",  
    totalOrderValue: 1,  
    customer_name: { $arrayElemAt: ["$customer_info.name", 0] }  
  }  
}  
});
```

In this example:

- The “\$group” stage groups orders by “customer_id” and calculates the total order value for each customer.
- The “\$lookup” stage performs a left outer join with the "customers" collection to retrieve customer information based on `customer_id`.
- The “\$project” stage reshapes the result to include “customer_id”, “totalOrderValue”, and “customer_name”.

This query yields a result that shows the total order value for each customer along with their names.

5

Data Modeling in MongoDB

Effective data modeling is essential for designing a database schema that aligns with your application's requirements and optimizes query performance.

5.1 Embedded vs. Referenced Documents

MongoDB's flexibility allows you to model your data in various ways. Two common approaches to consider when designing your schema are using embedded documents or referenced documents.

Embedded Documents

In this approach, you store related data within a single document. Embedded documents are useful for modeling one-to-one and one-to-many relationships, where the related data is relatively small and doesn't change frequently. Here's an example of embedding an address within a user document:

- **json**

```
{
```

```
"_id": 1,  
"name": "Alice",  
"email": "alice@example.com",  
"address": {  
  "street": "123 Main St",  
  "city": "Anytown",  
  "zipcode": "12345"  
}  
}
```

Referenced Documents

In this approach, you store a reference (usually an ObjectId) to related data in a separate collection. Referenced documents are useful for modeling many-to-one and many-to-many relationships, where the related data is large or may change frequently. Here's an example of referencing orders to products:

- json

// Products Collection

```
{  
  "_id": 101,  
  "name": "Laptop",  
  "price": 999.99
```



```
}
```

// Orders Collection

```
{
```

```
  "_id": 201,
```

```
  "user_id": 1,
```

```
  "products": [101, 102, 103]
```

```
}
```

The choice between embedding and referencing depends on your specific use case and the trade-offs between query performance, data consistency, and data duplication.

5.2 Data Normalization and De-Normalization

Data normalization and de-normalization are strategies for organizing your data to strike a balance between data integrity and query performance.

Data Normalization

Normalization is the process of organizing data in such a way that reduces data redundancy and ensures data consistency. It involves breaking down data into smaller, related pieces and storing them in separate collections. For example, you might store customer data in one collection and order data in another, using references to link

orders to customers. This approach minimizes data duplication and enforces data consistency.

- **json**

// Customers Collection

```
{  
  "_id": 1,  
  "name": "Alice",  
  "email": "alice@example.com"  
}
```

// Orders Collection

```
{  
  "_id": 101,  
  "customer_id": 1,  
  "total_amount": 200.0  
}
```

Data De-normalization

De-normalization, on the other hand, involves including redundant data within documents to optimize query performance. By duplicating certain data, you reduce the need for multiple queries

and JOIN operations. De-normalization is suitable for read-heavy workloads or situations where query performance is critical.

- **json**

// Users Collection with Embedded Orders

```
{
  "_id": 1,
  "name": "Alice",
  "email": "alice@example.com",
  "orders": [
    {
      "_id": 101,
      "total_amount": 200.0
    },
    {
      "_id": 102,
      "total_amount": 150.0
    }
  ]
}
```

The choice between normalization and de-normalization depends on your application's requirements. If you have a read-heavy workload

or need to optimize query performance, de-normalization can be a suitable choice. However, it may increase data storage requirements and complexity.

5.3 Modeling Relationship

MongoDB allows you to model various types of relationships between data.

One-to-One Relationship

For a one-to-one relationship, you can embed the related data within a document. For example, storing user profiles within the user document:

- **json**

// User Document with Embedded Profile

```
{
  "_id": 1,
  "username": "alice",
  "profile": {
    "first_name": "Alice",
    "last_name": "Smith",
    "age": 28
  }
}
```

```
}
```

One-to-Many Relationship

In a one-to-many relationship, you can embed an array of related documents within the parent document. For instance, storing comments within a blog post:

- json

// Blog Post Document with Embedded Comments

```
{  
  "_id": 101,  
  "title": "MongoDB Data Modeling",  
  "content": "..."  
  "comments": [  
    {  
      "_id": 1,  
      "user_id": 2,  
      "text": "Great article!"  
    },  
    {  
      "_id": 2,  
      "user_id": 3,
```

```
    "text": "Very informative."
  }
]
}
```

Many-to-Many Relationship

In a many-to-many relationship, you can use arrays of references to represent the relationships between documents. For example, modeling students and courses:

- json

// Students Collection

```
{
  "_id": 1,
  "name": "Alice",
  "courses": [101, 102]
}
```

// Courses Collection

```
{
  "_id": 101,
  "name": "Math"
```

```
}
```

In the **example** above, each student document references the courses they are enrolled in, and each course document is referenced by the students who are enrolled.

Example:

Let's consider a scenario where you're building an e-commerce platform. You have two main entities: "users" and "orders." Each user can place multiple orders. You can model this using embedded documents for orders within the user document:

- **json**

// Users Collection with Embedded Orders

```
{
  "_id": 1,
  "name": "Alice",
  "email": "alice@example.com",
  "orders": [
    {
      "_id": 101,
      "total_amount": 200.0
    },

```

```
{  
  "_id": 102,  
  "total_amount": 150.0  
}  
]  
}
```

In this example:

- Each user document contains an array of embedded order documents.
- Each order document includes details such as the order ID and total amount.

This schema simplifies queries for retrieving a user's orders, as you can directly access the orders within the user document. However, it can lead to data duplication if user information is repeated across multiple orders. The choice of whether to embed or reference orders would depend on factors like query patterns and the expected size of the orders.

6

Advanced MongoDB Features

In this module we will explore the features that enable you to work with specialized data types and perform advanced querying and analysis.

6.1 Geospatial Queries

MongoDB provides powerful support for geospatial data, allowing you to store and query data associated with geographical locations. This is especially useful for applications that require location-based services, such as mapping and “geofencing”.

To work with geospatial data in MongoDB, you typically store coordinates as part of your document and create a geospatial index on the relevant field. For example, consider a "locations" collection that stores information about various points of interest:

- json

// Sample "locations" collection

{

```
"_id": 1,  
"name": "Central Park",  
"location": {  
  "type": "Point",  
  "coordinates": [40.785091, -73.968285]  
}  
}
```

Here, the "location" field stores the latitude and longitude coordinates of Central Park. To perform geospatial queries, you can create a 2dsphere index:

- **javascript**

// Create a geospatial index on the "location" field

```
db.locations.createIndex({ location: "2dsphere" });
```

Now, you can execute geospatial queries like finding nearby locations:

- **javascript**

// Find locations within a specified radius (in meters) from a given point

```
db.locations.find({
```

```
location: {  
  $near: {  
    $geometry: {  
      type: "Point",  
      coordinates: [40.786, -73.964] // Example coordinates  
    },  
    $maxDistance: 500 // 500 meters radius  
  }  
}  
});
```

6.2 Text Search

MongoDB includes text search capabilities that allow you to perform full-text search operations on string fields. This feature is particularly useful for applications that need to search through large text datasets, such as content management systems or search engines.

To enable text search in MongoDB, you create a text index on one or more fields that you want to search. For example, consider a "products" collection with a "description" field:

- json

// Sample "products" collection

```
{  
  "_id": 1,  
  "name": "Laptop",  
  "description": "A powerful laptop for all your computing needs."  
}
```

To perform a text search, create a text index and use the “\$text” operator:

- **javascript**

// Create a text index on the "description" field

```
db.products.createIndex({ description: "text" });
```

Now, you can execute text searches:

- **javascript**

// Find products that contain the word "laptop" in the description

```
db.products.find({ $text: { $search: "laptop" } });
```

Text indexes support various features, including language-specific stemming, stop words, and relevance sorting, making them versatile for text-based search scenarios.

6.3 Full-Text Search with Text Indexes

MongoDB's text indexes support advanced text search capabilities, such as compound queries, phrase search, and fuzzy matching. These features enhance the precision and relevance of search results.

Compound Queries

You can combine multiple search terms using logical operators like “\$and”, “\$or”, and “\$not”. For example, to find products containing both "laptop" and "powerful" in the description:

- **javascript**

```
db.products.find({ $text: { $search: "laptop powerful" } });
```

Phrase Search

Use double quotes to search for an exact phrase. For example, to find products with the phrase "high-performance laptop":

- **javascript**

```
db.products.find({ $text: { $search: "\"high-performance laptop\"" } });
```

Fuzzy Matching

MongoDB supports fuzzy search, which finds similar terms to a specified query term. This is useful for handling misspellings or variations in user input:

- **javascript**

```
// Find products with terms similar to "laptop"
```

```
db.products.find({ $text: { $search: "laptap" } });
```

These advanced features enhance the precision and flexibility of full-text search in MongoDB.

6.4 Time Series Data with MongoDB

Many applications need to handle time-series data, which represents data points collected at specific time intervals. MongoDB offers features that make it suitable for time-series data storage and querying.

To work with time-series data, you typically structure your documents to include a timestamp along with the data point. For instance, consider a "sensor_data" collection that stores temperature readings:

- **json**

// Sample "sensor_data" collection

```
{  
  "_id": 1,  
  "timestamp": ISODate("2023-01-15T09:30:00Z"),  
  "temperature": 25.5  
}
```

To optimize queries for time-series data, you can create a compound index on the "timestamp" field and any other fields you frequently query or filter by:

- **javascript**

// Create a compound index on "timestamp" and "sensor_id" fields

```
db.sensor_data.createIndex({ timestamp: 1, sensor_id: 1 });
```

Now, you can perform efficient time-based queries, such as retrieving data points for a specific time range:

- **javascript**

// Find temperature readings for a specific sensor within a time range

```
db.sensor_data.find({  
  sensor_id: 101,
```

```
timestamp: {  
  $gte: ISODate("2023-01-15T00:00:00Z"),  
  $lt: ISODate("2023-01-16T00:00:00Z")  
}  
});
```

MongoDB's support for time-series data also extends to features like data retention policies, aggregation for summarizing time-series data, and support for various date and time operators.

Example:

Suppose you are building a location-based social network that allows users to post and search for places of interest. You want to support geospatial queries to find nearby places. Here's how you might model this in MongoDB:

- **json**

// "places" Collection with Geospatial Data

```
{  
  "_id": 1,  
  "name": "Central Park",  
  "location": {  
    "type": "Point",  
    "coordinates": [40.785091, -73.968285]
```



```
},  
"category": "Park"  
}
```

To find places near a user's location, you can execute a geospatial query with a `\$near` operator:

- javascript

// Find places near the user's location

```
db.places.find(  
  location: {  
    $near: {  
      $geometry: {  
        type: "Point",  
        coordinates: [40.786, -73.964] // Example user's coordinates  
      },  
      $maxDistance: 500 // 500 meters radius  
    }  
  }  
});
```

This query returns a list of nearby places within a 500-meter radius of the user's location.

7

MongoDB Atlas and Cloud Deployment

In this module we will explore the features of MongoDB Atlas, MongoDB's official cloud-based database service. MongoDB Atlas provides a convenient way to host, manage, and scale MongoDB databases in a cloud environment, making it an essential topic for modern application development.

7.1 Introduction to MongoDB Atlas

MongoDB Atlas is a fully managed database service that allows you to deploy, manage, and scale MongoDB databases in the cloud. It offers several advantages for developers and organizations:

Automated Management: MongoDB Atlas handles routine database management tasks, such as hardware provisioning, setup, and configuration, leaving you free to focus on application development.

Scalability: You can easily scale your MongoDB Atlas clusters up or down to meet the demands of your application, ensuring optimal performance.

Security: MongoDB Atlas provides robust security features, including encryption, authentication, and network isolation, to protect your data.

Backup and Recovery: Automated backups and point-in-time recovery options help you safeguard your data against loss or corruption.

Monitoring and Insights: MongoDB Atlas offers monitoring and performance optimization tools to help you identify and address potential issues in your database.

Global Deployment: You can deploy MongoDB Atlas clusters in multiple regions to reduce latency and provide a better experience for users around the world.

7.2 Creating and Managing Clusters

A cluster in MongoDB Atlas represents a group of MongoDB servers that work together to store your data. Clusters come in various configurations to accommodate different workloads and performance requirements. To create and manage clusters in MongoDB Atlas, follow these steps:

Step 1: Create an Atlas Account

If you don't already have an account, sign up for MongoDB Atlas at <https://www.mongodb.com/cloud/atlas>.

Step 2: Create a New Project

Projects in MongoDB Atlas help you organize your database resources. Create a new project and give it a descriptive name.

Step 3: Create a Cluster

Within your project, click the "Clusters" tab and then click the "Build a New Cluster" button. You'll be prompted to select various configuration options for your cluster, including:

- Cloud provider (e.g., AWS, Azure, Google Cloud)
- Region and availability zone
- Cluster tier (e.g., M10, M30, M60, etc.)
- Additional features (e.g., backup, monitoring)

After configuring your cluster, click the "Create Cluster" button.

Step 4: Configure Cluster Settings

Once your cluster is created, you can configure additional settings such as network access, security, and data storage.

Network Access: Specify which IP addresses or IP ranges are allowed to connect to your cluster. You can whitelist IPs for your application servers.

Security: Configure authentication and authorization settings to secure access to your database.

Data Storage: Adjust storage settings, including enabling automated backups and choosing a backup retention policy.

Step 5: Connect to Your Cluster

MongoDB Atlas provides connection strings that you can use to connect your application to the cluster. These strings include authentication details and other connection parameters. You can choose between different driver-specific connection strings.

Here's an example of a connection string for connecting a Node.js application to a MongoDB Atlas cluster:

- **javascript**

```
const MongoClient = require("mongodb").MongoClient;
```

```
const uri =
"mongodb+srv://<username>:<password>@clustername.mongodb.net/test?retryWrites=true&w=majority";

MongoClient.connect(uri, (err, client) => {
  if (err) {
    console.error("Error connecting to MongoDB:", err);
    return;
  }

  const db = client.db("mydatabase");

  // Your database operations here

  client.close();
});
```

7.3 Deploying MongoDB in the Cloud

MongoDB Atlas simplifies the process of deploying MongoDB databases in the cloud. With a few clicks, you can provision and configure MongoDB clusters in your preferred cloud provider's infrastructure.

Here's a step-by-step guide to deploying MongoDB in the cloud using MongoDB Atlas:

Step 1: Select Cloud Provider

MongoDB Atlas supports major cloud providers, including AWS, Azure, and Google Cloud Platform (GCP). Choose the provider that suits your needs.

Step 2: Choose Region

Select the region or data center location where you want to deploy your MongoDB cluster. Consider factors like latency and data residency requirements when making your choice.

Step 3: Configure Cluster Tier

Choose the appropriate cluster tier based on your application's performance requirements and budget. MongoDB Atlas offers various cluster tiers with different compute and storage capacities.

Step 4: Set Additional Options

Configure additional options for your cluster, such as backup settings, maintenance windows, and auto-scaling options.

Step 5: Secure Your Cluster

Set up security features like network access control, authentication, and encryption to protect your data.

Step 6: Review and Create

Review your cluster configuration, including pricing details, before creating the cluster. Once you're satisfied, click the "Create Cluster" button.

Step 7: Connect to Your Cluster

MongoDB Atlas provides connection strings that you can use to connect your application to the newly created cluster. Follow the provided guidelines to establish a connection.

Example Code:

Below is an example of Node.js code to connect to a MongoDB Atlas cluster:

- javascript

```
const MongoClient = require("mongodb").MongoClient;
```

```
const uri =  
"mongodb+srv://<username>:<password>@clustername.mongodb.net/test?retryWrites=true&w=majority";
```

```
MongoClient.connect(uri, (err, client) => {
```

```
  if (err) {
```

```
    console.error("Error connecting to MongoDB:", err);
```



```
    return;  
}  
  
const db = client.db("mydatabase");  
  
// Your database operations here  
  
client.close();  
});
```

In this example:

- Replace "<username>" and "<password>" with your MongoDB Atlas credentials.
- Replace ""clustername"" with the actual name of your MongoDB Atlas cluster.
- You can specify the database you want to connect to (e.g., "mydatabase") in the "db" variable.