

Modules

Module – 1 Introduction to MongoDB

- 1.1 Overview of MongoDB**
- 1.2 NoSQL vs. SQL databases**
- 1.3 Advantages of using MongoDB**
- 1.4 Installing and setting up MongoDB**

Module – 2 MongoDB Data Model

- 2.1 Understanding Documents and Collections**
- 2.2 BSON Data Format**
- 2.3 Document-Oriented Data Model**
- 2.4 MongoDB Schema Design Best Practices**

Module – 3 CRUD Operations in MongoDB

- 3.1 Creating documents**
- 3.2 Reading documents**
- 3.3 Updating documents**
- 3.4 Deleting documents**
- 3.5 Querying data with MongoDB**

Module – 4 Indexing and Query Optimization

- 4.1 Importance of indexing**
- 4.2 Creating and Managing Indexes**
- 4.3 Query Optimization Techniques**
- 4.4 Using the Aggregation Framework**

Module – 5 Data Modeling in MongoDB

- 5.1 Embedded vs. Referenced documents**
- 5.2 Data normalization and de-normalization**
- 5.3 Modeling Relationship**

Module – 6 Advanced MongoDB Features

- 6.1 Geospatial Queries**
- 6.2 Text search**
- 6.3 Full-Text Search with Text Indexes**
- 6.4 Time-Series Data with MongoDB**

Module – 7 MongoDB Atlas and Cloud Deployment

- 7.1 Introduction to MongoDB Atlas**
- 7.2 Creating and Managing Clusters**
- 7.3 Deploying MongoDB in the Cloud**

Module – 8 Security and Authentication

- 8.1 Securing MongoDB**
- 8.2 User Authentication and Authorization**
- 8.3 Role-Based Access Control (RBAC)**

Module – 9 Backup and Recovery

- 9.1 Backup Strategies**
- 9.2 Restoring Data**
- 9.3 Disaster Recovery Planning**

Module – 10 MongoDB Aggregation Pipeline

- 10.1 Aggregation Concepts**
- 10.2 Using Pipeline Stages**
- 10.3 Custom Aggregation Expressions**

Module – 11 MongoDB Drivers and APIs

- 11.1 MongoDB Drivers**
- 11.2 Programming Language of your Choice**

Module – 12 MongoDB Performance Tuning

- 12.1 Profiling and Monitoring**
- 12.2 Query Performance Optimization**
- 12.3 Hardware and Resource Considerations**

Module – 13 Replication and Sharding

- 13.1 Replication for High Availability**
- 13.2 Sharding for Horizontal Scaling**
- 13.3 Configuring Replica Sets and Sharded Clusters**

Module – 14 Working with GridFS

- 14.1 Storing Large files in MongoDB**
- 14.2 Using GridFS for file Management**

8

Security and Authentication

In this module we will explore the aspects of securing MongoDB databases, implementing user authentication and authorization, and leveraging role-based access control (RBAC) to manage permissions. Security is paramount in any database system, and MongoDB provides robust features to protect your data from unauthorized access and potential threats.

8.1 Securing MongoDB

Securing MongoDB involves a combination of measures to protect your database from unauthorized access, data breaches, and potential security vulnerabilities. Some essential security practices include:

- 1. Firewall Configuration:** MongoDB should be configured to only accept connections from trusted IP addresses or networks. You can use the IP Whitelist feature in MongoDB Atlas or configure network interfaces in on-premises installations.

2. Enable Authentication: MongoDB should always require authentication. This means that users must provide valid credentials (username and password) to access the database.

3. Use Encryption: Data in transit should be encrypted using TLS/SSL. MongoDB supports encrypted connections between clients and the database server.

4. Patch and Update: Keep MongoDB up to date with the latest security patches and updates to mitigate vulnerabilities.

5. Least Privilege Principle: Grant only the minimum necessary privileges to users and applications. Avoid giving overly broad permissions.

6. Audit and Logging: Enable auditing and logging to track and monitor database activities for security incidents and compliance.

7. Secure Configuration: Configure MongoDB with security in mind, using secure settings and following best practices.

8.2 User Authentication and Authorization

User authentication and authorization are fundamental components of database security. Authentication ensures that users are who they

claim to be, while authorization controls what actions and data they can access.

Authentication in MongoDB

MongoDB supports various authentication methods, including username/password, LDAP (Lightweight Directory Access Protocol), and x.509 certificates. To enable authentication, you need to create user accounts and specify authentication mechanisms.

Here's an example of creating a user with username and password authentication:

- **javascript**

// Connect to the MongoDB server

use admin

```
db.createUser({  
  user: "myuser",  
  pwd: "mypassword",  
  roles: [{ role: "readWrite", db: "mydatabase" }]  
});
```

In this example:

- We first switch to the "admin" database where user accounts are typically created.

- We use the ``createUser`` method to create a new user with the username "myuser" and password "mypassword."
- We grant the user the "readWrite" role for the "mydatabase" database, allowing them to read and write data in that database.

Authorization in MongoDB

Once users are authenticated, you can use MongoDB's role-based access control (RBAC) to define their permissions. RBAC allows you to specify roles with specific privileges and assign those roles to users.

Here's an example of creating a custom role and assigning it to a user:

- **javascript**

// Create a custom role with read-only permissions on a specific collection

```
db.createRole({
  role: "customReadOnly",
  privileges: [
    {
      resource: { db: "mydatabase", collection: "mycollection" },
      actions: ["find"]
    }
  ]
})
```

```
}  
],  
roles: []  
});
```

// Create a user and assign the custom role

```
db.createUser({  
  user: "customUser",  
  pwd: "userpassword",  
  roles: [{ role: "customReadOnly", db: "mydatabase" }]  
});
```

In this example:

- We create a custom role called "customReadOnly" with privileges that allow users to perform the "find" action on the "mycollection" collection within the "mydatabase" database.
- We then create a user named "customUser" with the password "userpassword" and assign them the "customReadOnly" role, giving them read-only access to the specified collection.

8.3 Role-Based Access Control (RBAC)

MongoDB's RBAC system allows you to manage permissions and access control at a granular level. Roles can be predefined (built-in

roles) or custom (user-defined roles). MongoDB provides several built-in roles, such as "read", "readWrite", and "dbAdmin", which offer predefined sets of permissions.

Here's an example of creating a user-defined role that can perform administrative actions on a specific database:

- **javascript**

// Create a custom role with administrative privileges on a specific database

```
db.createRole({
  role: "customDbAdmin",
  privileges: [
    {
      resource: { db: "mydatabase", collection: "" },
      actions: ["listCollections"]
    },
    {
      resource: { db: "mydatabase", collection: "" },
      actions: ["createCollection"]
    }
  ],
  roles: []
})
```

```
});
```

In this example:

We create a custom role called "customDbAdmin" with privileges that allow users to list collections and create collections within the "mydatabase" database.

Once the role is defined, you can assign it to specific users or grant it to other roles.

Example:

Here's an example of connecting to a secured MongoDB database using authentication in Node.js:

- javascript

```
const MongoClient = require("mongodb").MongoClient;
```

```
const uri = "mongodb://myuser:mypassword@mongodb-server/mydatabase";
```

```
MongoClient.connect(uri, (err, client) => {
```

```
  if (err) {
```

```
    console.error("Error connecting to MongoDB:", err);
```

```
    return;  
}  
  
const db = client.db("mydatabase");  
  
// Your database operations here  
  
client.close();  
});
```

In this example:

- Replace `"myuser"` and `"mypassword"` with the appropriate username and password for authentication.
- Replace `"mongodb-server"` with the hostname or IP address of your MongoDB server.
- Specify the database you want to connect to using `"mydatabase"`.

9

Backup and Recovery

In this module we will explore the aspects of backup and recovery in MongoDB. Effective backup and recovery strategies are essential for safeguarding your data, ensuring data availability, and mitigating the impact of data loss or system failures. This module covers various backup strategies, data restoration techniques, and disaster recovery planning for MongoDB.

9.1 Backup Strategies

Backup strategies in MongoDB involve creating copies of your data and storing them securely to prevent data loss due to various factors, such as hardware failures, accidental deletions, or system errors. MongoDB provides several methods and tools for performing backups:

1. `mongodump` and `mongorestore`

MongoDB includes the `mongodump` and `mongorestore` utilities, which allow you to create and restore backups at the database or collection level. These utilities generate BSON (Binary JSON) dump

files that capture the data and indexes in your database.

Backup Example:

To create a backup of a specific database using “mongodump”, you can run the following command:

- **shell**

```
mongodump --host <hostname> --port <port> --db  
<database_name> --out <backup_directory>
```

- Replace “<hostname>” and “<port>” with the MongoDB server's hostname and port.
- Specify the “<database_name>” you want to back up.
- Provide the “<backup_directory>” where the dump files will be saved.

Restore Example:

To restore a database from a “mongodump” backup, use the “mongorestore” command:

- **shell**

```
mongorestore --host <hostname> --port <port> --db  
<database_name> <backup_directory>/<database_name>
```

- Replace “<hostname>” and “<port>” with the MongoDB server's hostname and port.
- Specify the “<database_name>” to restore.
- Provide the path to the backup directory where the dump files are stored.

2. Filesystem Snapshots

Another backup approach is to take filesystem snapshots at the storage level. This method involves creating point-in-time snapshots of the entire MongoDB data directory. While this approach is efficient, it requires support from your storage infrastructure and may not be suitable for all environments.

3. Cloud Backup Services

MongoDB Atlas, MongoDB's official cloud service, provides automated backup solutions. MongoDB Atlas offers daily snapshots of your data, which you can restore from directly using the Atlas interface. This approach simplifies the backup process and ensures data availability in the cloud.

4. Third-Party Backup Solutions

There are third-party backup solutions and services available for MongoDB that offer additional features and customization options. These solutions may be suitable for enterprises with specific backup and recovery requirements.

9.2 Restoring Data

Data restoration is the process of recovering data from backups when needed. MongoDB provides various methods for restoring data, depending on the backup strategy used:

1. mongorestore

As mentioned earlier, you can use the “mongorestore” utility to restore data from “mongodump” backups. This utility can restore data at the database or collection level.

Example:

To restore a specific database from a “mongodump” backup, you can use the following command:

- shell

```
mongorestore --host <hostname> --port <port> --db  
<database_name> <backup_directory>/<database_name>
```

2. Atlas Restore

If you are using MongoDB Atlas, you can restore data directly from the Atlas interface. Atlas provides a user-friendly interface for selecting and restoring snapshots of your data. You can choose the specific point-in-time snapshot you want to restore, and Atlas will handle the process.

3. Filesystem Snapshots

For backups created using filesystem snapshots, you can restore data by reverting the MongoDB data directory to a specific snapshot. This process typically involves working with your storage infrastructure and file system snapshots.

9.3 Disaster Recovery Planning

Disaster recovery planning is an essential part of ensuring the resilience of your MongoDB deployments. It involves preparing for unforeseen events that can lead to data loss or system downtime, such as hardware failures, natural disasters, or cyberattacks. Here are key considerations for disaster recovery planning in MongoDB:

- 1. Identify Critical Data:** Determine which data is critical for your organization and prioritize its backup and recovery.
- 2. Regular Backups:** Implement regular backup schedules to ensure that data is continuously protected.
- 3. Offsite Backup Storage:** Store backup copies in geographically separate locations to protect against local disasters.

4. Test Restores: Regularly test the restoration process to ensure that backups are reliable and can be successfully restored when needed.

5. Disaster Recovery Plan: Develop a comprehensive disaster recovery plan that outlines procedures for various disaster scenarios, including data restoration and system recovery.

6. Monitoring and Alerts: Implement monitoring and alerting systems to detect issues early and take preventive actions.

7. Backup Retention Policies: Define backup retention policies to manage how long backups are retained and when older backups can be deleted.

Example Code:

Here's an example of creating a backup using `mongodump` and then restoring the backup using `mongorestore`:

Backup:

- shell

Create a backup using mongodump

```
mongodump --host <hostname> --port <port> --db mydatabase --out /backup
```

Restore:

- shell

Restore the backup using mongorestore

```
mongorestore --host <hostname> --port <port> --db mydatabase  
/backup/mydatabase
```

In these commands:

- Replace “<hostname>” and “<port>” with the MongoDB server's hostname and port.
- Use the “mongodump” command to create a backup in the “/backup” directory.
- Use the “mongorestore” command to restore the backup to the "mydatabase" database.

10

MongoDB Aggregation Pipeline

In this module we will explore the MongoDB Aggregation Pipeline, a powerful tool for data transformation and analysis within MongoDB. The Aggregation Pipeline allows you to perform complex operations on your data, including filtering, grouping, sorting, and calculating aggregations.

10.1 Aggregation Concepts

Aggregation in MongoDB refers to the process of transforming and summarizing data within a collection. It allows you to analyze and manipulate data to extract meaningful insights. Aggregation operations can involve multiple stages, and the Aggregation Pipeline provides a framework for defining these stages.

Key aggregation concepts include:

- 1. Pipeline:** The aggregation pipeline is a sequence of stages, where each stage represents an operation to be performed on the data.

Data flows through these stages sequentially, and each stage produces intermediate results that feed into the next stage.

2. Stage: A stage is a specific operation or transformation applied to the data. Common stages include “\$match”, “\$group”, “\$sort”, and “\$project”, among others.

3. Document Transformation: Aggregation can transform documents by filtering, reshaping, and computing new fields. This allows you to tailor the output to your specific requirements.

4. Grouping: The “\$group” stage is used to group documents by specified fields and perform aggregation operations within each group. Aggregation functions like “\$sum”, “\$avg”, “\$min”, and “\$max” can be applied to grouped data.

5. Sorting: The “\$sort” stage allows you to sort the aggregated results based on one or more fields in ascending or descending order.

6. Expression Operators: MongoDB provides a wide range of expression operators that can be used in aggregation stages to perform arithmetic, logical, and comparison operations on fields.

10.2 Using Pipeline Stages

The MongoDB Aggregation Pipeline consists of multiple stages, and each stage performs a specific operation on the data. Here are some commonly used aggregation pipeline stages:

1. \$match: This stage filters documents based on specified criteria, allowing you to select a subset of documents to include in the aggregation.

Example:

- **javascript**

```
db.sales.aggregate([  
  { $match: { date: { $gte: ISODate("2022-01-01"), $lte:  
    ISODate("2022-12-31") } } }  
]);
```

In this example, the “\$match” stage filters sales documents for the year 2022.

2. \$group: The “\$group” stage groups documents by one or more fields and calculates aggregations within each group.

Example:

- **javascript**

```
db.sales.aggregate([
```

```
{ $group: { _id: "$product", totalSales: { $sum: "$quantity" } } }  
]);
```

This stage groups sales by product and calculates the total quantity sold for each product.

3. \$project: The “\$project” stage reshapes documents by including or excluding fields, creating new fields, or applying expressions to existing fields.

Example:

- **javascript**

```
db.sales.aggregate([  
  { $project: { _id: 0, product: 1, revenue: { $multiply: ["$price",  
"$quantity"] } } }  
]);
```

Here, the ` \$project ` stage calculates the revenue for each sale and includes only the "product" and "revenue" fields in the output.

4. \$sort: The “\$sort” stage sorts the documents based on specified fields and sort order.

Example:

- **javascript**

```
db.sales.aggregate([  
  { $sort: { revenue: -1 } }  
]);
```

This stage sorts sales documents in descending order of revenue.

5. \$limit and \$skip: These stages allow you to limit the number of documents returned in the result set and skip a specified number of documents.

Example:

- **javascript**

```
db.sales.aggregate([  
  { $sort: { revenue: -1 } },  
  { $limit: 5 },  
  { $skip: 2 }  
]);
```

This sequence first sort's sales document by revenue, then limits the result to the top 5, and finally skips the first 2.

10.3 Custom Aggregation Expressions

In MongoDB Aggregation, you can use custom aggregation expressions to perform calculations and transformations on your data. These expressions are built using aggregation operators and can be used within various stages to manipulate documents.

Examples of custom aggregation expressions include:

Arithmetic Operations:

- **javascript**

```
db.sales.aggregate([  
  {  
    $project: {  
      total: {  
        $add: ["$price", "$tax"]  
      }  
    }  
  }  
]);
```

In this example, the “\$add” operator calculates the sum of the "price" and "tax" fields.

Logical Operations

- javascript

```
db.students.aggregate([
  {
    $project: {
      passed: {
        $eq: ["$score", { $literal: 100 }]
      }
    }
  }
]);
```

Here, the “\$eq” operator checks if the "score" field is equal to 100.

String Manipulation

- javascript

```
db.contacts.aggregate([
  {
    $project: {
      fullName: {
```

```
    $concat: ["$firstName", " ", "$lastName"]
  }
}
}
]);
```

The “\$concat” operator concatenates the "firstName" and "lastName" fields to create a "fullName" field.

Conditional Expressions

- javascript

```
db.orders.aggregate([
{
  $project: {
    status: {
      $cond: {
        if: { $eq: ["$shipped", true] },
        then: "Shipped",
        else: "Pending"
      }
    }
  }
}]
```

```
}  
}  
]);
```

The “\$cond” operator applies a conditional expression to determine the "status" field value based on the "shipped" field.

Date Operations

- javascript

```
db.events.aggregate([  
  {  
    $project: {  
      formattedDate: {  
        $dateToString: { format: "%Y-%m-%d", date: "$eventDate" }  
      }  
    }  
  }  
]);
```

The “\$dateToString” operator formats the "eventDate" field as a string in the specified format.

11

MongoDB Drivers and APIs

In this module we will explore the MongoDB drivers and APIs, which are essential components for interacting with MongoDB databases using various programming languages. MongoDB offers official drivers and community-supported libraries for many programming languages, making it accessible and versatile for developers.

11.1 MongoDB Drivers for various Languages

MongoDB drivers are software libraries or modules that enable developers to connect to and interact with MongoDB databases using specific programming languages. MongoDB provides official, well-maintained drivers for popular programming languages like Python, Node.js, Java, C#, and more. These drivers offer comprehensive functionality and are continuously updated to support the latest MongoDB features.

Here are some of the official MongoDB drivers for popular programming languages:

Python: PyMongo is the official MongoDB driver for Python, allowing Python developers to work seamlessly with MongoDB databases.

Node.js: The MongoDB Node.js driver enables Node.js developers to build scalable, high-performance applications with MongoDB.

Java: MongoDB provides a Java driver for Java-based applications, ensuring compatibility and performance.

C#: The official C# driver (MongoDB .NET Driver) enables developers using .NET languages like C# to interact with MongoDB.

Ruby: Ruby developers can use the Ruby driver (MongoDB Ruby Driver) for MongoDB integration.

Go: The Go programming language has an official MongoDB driver known as the MongoDB Go Driver.

PHP: PHP developers can use the MongoDB PHP driver for building web applications with MongoDB.

Scala: The Scala driver (MongoDB Scala Driver) is available for Scala applications.

Swift: For iOS and macOS app development, the official MongoDB Swift driver provides seamless integration.

Kotlin: Kotlin developers can use the official Kotlin driver (KMongo) for MongoDB.

Rust: The Rust programming language has the official MongoDB Rust driver for building efficient and safe applications.

11.2 Programming Language of Your Choice

Working with MongoDB Using a Programming Language of Your Choice

To work with MongoDB using a programming language of your choice, you need to follow these general steps:

1. Install the MongoDB Driver

Begin by installing the MongoDB driver for your chosen programming language. You can usually install the driver using a package manager or include it as a dependency in your project's configuration.

2. Import or Include the Driver

Import or include the MongoDB driver in your code to access its features and functions. This typically involves using the appropriate import statements or directives.

3. Establish a Connection

Connect to your MongoDB database by specifying the connection details, such as the hostname, port, and authentication credentials. Most MongoDB drivers provide connection pooling for efficient and reusable connections.

4. Perform CRUD Operations

Use the driver's API to perform CRUD (Create, Read, Update, Delete) operations on your MongoDB data. You can insert documents, query data, update records, and delete documents as needed.

5. Handle Errors and Exceptions

Be prepared to handle errors and exceptions that may occur during database interactions. This includes handling network errors, authentication failures, and data validation errors.

6. Close the Connection

After you have finished working with the database, remember to close the database connection to release resources properly.

Example Code (Python - PyMongo):

Here's an example of using the PyMongo driver to connect to a MongoDB database and perform basic operations in Python:

- **python**

```
from pymongo import MongoClient
```

```
# Establish a connection to the MongoDB server
```

```
client = MongoClient("mongodb://localhost:27017/")
```

```
# Access a specific database
```

```
db = client["mydatabase"]
```

```
# Access a collection within the database
```

```
collection = db["mycollection"]
```

```
# Insert a document
```



```
data = {"name": "John", "age": 30, "city": "New York"}  
inserted_id = collection.insert_one(data).inserted_id
```

Query for documents

```
result = collection.find({"age": {"$gte": 25}})  
  
for document in result:  
    print(document)
```

Update a document

```
collection.update_one({"_id": inserted_id}, {"$set": {"city": "San  
Francisco"}})
```

Delete a document

```
collection.delete_one({"_id": inserted_id})
```

Close the MongoDB connection

```
client.close()
```

In this Python example:

- We import the MongoClient class from PyMongo and establish a connection to a MongoDB server running locally.
- We access a specific database ("mydatabase") and a collection ("mycollection") within that database.

- We insert a document, query for documents matching a condition (age greater than or equal to 25), update a document, and delete a document.
- Finally, we close the MongoDB connection using the `close()` method.

This code demonstrates the basic operations you can perform with a MongoDB driver in Python. Similar operations can be performed with MongoDB drivers for other programming languages, tailored to the language's syntax and conventions.