

# 12

## MongoDB Performance Tuning

In this module we will explore the MongoDB performance tuning, a critical aspect of database administration and application development. Optimizing MongoDB performance ensures that your database operates efficiently, delivers fast query response times, and can handle increased workloads.

### 12.1 Profiling and Monitoring

Effective profiling and monitoring are fundamental to identifying and addressing performance issues in MongoDB. Profiling involves collecting data about database operations, while monitoring entails tracking the overall health and performance of the MongoDB deployment.

Key profiling and monitoring concepts include:

#### 1. Database Profiling

MongoDB allows you to enable database profiling to collect data on slow-running queries and operations. Profiling data can help identify bottlenecks and areas for optimization.

### **Example of enabling profiling:**

- **javascript**

```
db.setProfilingLevel(1, { slowms: 100 });
```

In this example, profiling is enabled at level 1, and queries taking longer than 100 milliseconds are logged.

## **2. Monitoring Tools**

MongoDB provides tools like the MongoDB Atlas Performance Advisor and third-party monitoring solutions to help you visualize and analyze the performance of your MongoDB deployment. These tools offer insights into resource utilization, query execution times, and other performance-related metrics.

## **3. Query Profiling**

Profiling can be used to identify slow queries by examining the “system.profile” collection. This collection stores profiling data, including the query's execution time and other relevant information.

### **Example of querying profiling data:**

- **javascript**

```
db.system.profile.find({ millis: { $gt: 100 } }).sort({ ts: -1 }).pretty();
```

This query retrieves profiling data for queries taking longer than 100 milliseconds and sorts the results by timestamp.

## 12.2 Query Performance Optimization

Optimizing query performance is crucial for delivering fast response times and improving the overall efficiency of MongoDB. Effective query optimization involves various strategies and techniques.

### 1. Indexing

Properly indexing your MongoDB collections can significantly improve query performance. Indexes allow MongoDB to quickly locate and retrieve documents that match specific criteria.

#### Example of creating an index:

- **javascript**

```
db.mycollection.createIndex({ field1: 1, field2: -1 });
```

This command creates a compound index on "field1" in ascending order and "field2" in descending order.

### 2. Query Structure

Optimize query structures to minimize the data returned by queries. Use the "select" option to specify the fields to return and

filter results using query operators like ``$eq``, ``$gt``, ``$lt``, and ``$in`` to narrow down the result set.

### **Example of an optimized query:**

#### **- javascript**

```
db.mycollection.find({ status: "active" }, { name: 1, date: 1
}).limit(10);
```

This query fetches only the "name" and "date" fields for documents with a "status" of "active" and limits the result set to 10 documents.

### **3. Avoid Large Result Sets**

When dealing with large collections, use pagination to retrieve results in smaller batches rather than fetching all documents at once. The “skip()” and “limit()” methods can help with pagination.

### **Example of pagination:**

#### **- javascript**

```
const pageSize = 10;
```

```
const pageNumber = 1;
```

```
const skipAmount = (pageNumber - 1) * pageSize;
```

```
db.mycollection.find({}).skip(skipAmount).limit(pageSize);
```

This query retrieves the first page of results with a page size of 10.

#### **4. Use Covered Queries**

Covered queries occur when all the fields needed for a query are present in an index, eliminating the need to access the actual documents. Covered queries are generally faster and more efficient.

##### **Example of a covered query:**

- **javascript**

```
db.mycollection.find({ field1: "value1" }, { _id: 0, field2: 1 });
```

In this query, "field2" is projected from the index, making it a covered query.

## **12.3 Hardware and Resource Consideration**

MongoDB performance can also be influenced by hardware and resource allocation. Properly configuring and managing hardware resources is crucial for optimal performance.

### **1. Memory (RAM)**

MongoDB benefits significantly from having sufficient RAM to store frequently accessed data and indexes. Ensure that the working set (the portion of data most frequently accessed) fits in RAM to avoid frequent disk I/O.

## **2. Disk Speed and Storage**

High-performance disks, such as SSDs, can greatly improve read and write operations. Monitor disk usage and consider horizontal scaling if storage becomes a bottleneck.

## **3. CPU Cores**

MongoDB can leverage multiple CPU cores for parallel processing. Ensure that your server has an adequate number of CPU cores to handle concurrent requests.

## **4. Network Throughput**

Network speed and throughput can affect data transfer rates between MongoDB servers and client applications. High-speed networks can reduce latency.

## **5. MongoDB Configuration**

MongoDB configuration options, such as the storage engine, write concern, and read preference, can impact performance. Review and optimize these settings based on your specific use case.

# 13

## Replication and Sharding in MongoDB

In this module we will explore two essential MongoDB features: replication and sharding. Replication ensures high availability and data redundancy, while sharding enables horizontal scaling for large datasets.

### 13.1 Replication of High Availability

Replication is the process of creating and maintaining multiple copies (replicas) of your data on different servers. It offers several benefits, including high availability, data redundancy, and fault tolerance. MongoDB implements replication through replica sets, which consist of multiple MongoDB instances.

Key concepts related to replication in MongoDB include:

- 1. Primary and Secondaries:** In a replica set, one member serves as the primary node, handling all write operations and becoming the authoritative source of data. The other members are secondaries

that replicate data from the primary. If the primary fails, one of the secondaries can be automatically elected as the new primary.

**2. Automatic Failover:** MongoDB replica sets provide automatic failover, ensuring that if the primary node becomes unavailable, one of the secondaries is automatically promoted to primary status. This minimizes downtime and ensures data availability.

**3. Read Scaling:** Read operations can be distributed across secondary nodes, allowing you to scale read-intensive workloads horizontally.

**4. Data Redundancy:** Data is replicated to multiple nodes, providing redundancy and reducing the risk of data loss due to hardware failures.

## **Configuring Replica Sets**

To configure a MongoDB replica set, you'll need to follow these general steps:

### **1. Initialize the Replica Set**

Initialize the replica set by connecting to one of the MongoDB nodes and running the “rs.initiate()” command.



### **Example:**

- **javascript**

```
rs.initiate({  
  _id: "myreplicaset",  
  members: [  
    { _id: 0, host: "mongo1:27017" },  
    { _id: 1, host: "mongo2:27017" },  
    { _id: 2, host: "mongo3:27017" }  
  ]  
});
```

In this example, a replica set named "myreplicaset" is initiated with three members.

## **2. Add Members**

You can add additional members to the replica set to increase redundancy and distribute read operations.

### **Example of adding a member:**

- **javascript**

```
rs.add("mongo4:27017");
```

This command adds a new member to the replica set.

### 3. Configure Read Preferences

Configure your application to use appropriate read preferences to route read operations to secondary nodes for read scaling.

**Example** of setting a read preference in the MongoDB Node.js driver:

- **javascript**

```
const MongoClient = require("mongodb").MongoClient;
```

```
const uri = "mongodb://mongo1:27017,mongo2:27017,mongo3:27017/?replicaSet=myreplicaset";
```

```
const client = new MongoClient(uri, { readPreference: "secondary" });
```

In this example, read operations will be distributed to secondary nodes.

## 13.2 Sharding for Horizontal Scaling

Sharding is a technique used to horizontally partition large datasets across multiple servers, or shards, to achieve horizontal scaling. MongoDB implements sharding through sharded clusters, which consist of multiple shard servers and configuration servers.

Key concepts related to sharding in MongoDB include:

- 1. Shard Key:** The shard key is a field in the documents that determines how data is distributed across shards. Choosing an appropriate shard key is critical for evenly distributing data and optimizing query performance.
- 2. Chunks:** Data is divided into smaller units called chunks. Each chunk is associated with a specific range of shard key values and is stored on a particular shard.
- 3. Balancing:** MongoDB automatically balances data distribution across shards by migrating chunks between shards as needed. This ensures that each shard has a roughly equal amount of data.
- 4. Config Servers:** Config servers store metadata about sharded clusters, including information about the shard key and chunk ranges.

## 13.3 Configuring Replica Sets - Sharded Cluster

To configure a MongoDB sharded cluster, you'll need to follow these general steps:

### 1. Initialize Config Servers

Initialize the config servers by starting multiple MongoDB instances as config servers and specifying the `--configsvr` option.

#### Example:

- **shell**

```
mongod --configsvr --replSet configReplSet --bind_ip localhost --port 27019
```

In this example, a config server is started with replication set "configReplSet."

### 2. Initialize Shards

Start multiple MongoDB instances to serve as shard servers. Each shard server should be started with the `--shardsvr` option.

#### Example:

- **shell**

```
mongod --shardsvr --replSet shard1ReplSet --bind_ip localhost --port 27018
```

This command starts a shard server with replication set "shard1ReplSet."

### **3. Initialize Mongos Routers**

Start Mongos routers, which are query routers that route client requests to the appropriate shard. Mongos instances should be aware of the config servers and shards.

**Example:**

- **shell**

```
mongos --configdb configReplSet/localhost:27019 --bind_ip localhost --port 27017
```

In this example, a Mongos router is started with knowledge of the config servers.

### **4. Enable Sharding**

Enable sharding for a specific database by connecting to a Mongos instance and running the `sh.enableSharding()` command.

**Example:**

**- javascript**

```
use mydatabase
```

```
db.createCollection("mycollection")
```

```
sh.enableSharding("
```

```
mydatabase")
```

```
  sh.shardCollection("mydatabase.mycollection", { shardKeyField: 1  
  })
```

This code enables sharding for the "mydatabase" database and specifies the shard key field.

## **5. Balancing Data**

MongoDB will automatically balance data across shards by moving chunks between them. No manual intervention is required.

By configuring sharded clusters, you can horizontally scale your MongoDB deployment to handle large datasets and high workloads efficiently.

# 14

## Working with GridFS

In this module we will explore GridFS, a specification within MongoDB for storing and retrieving large files and binary data. GridFS is particularly useful for handling files that exceed MongoDB's document size limit of 16 MB.

### 14.1 Storing Large files in MongoDB

MongoDB is designed to store structured JSON-like documents, and while it's great for most types of data, it has a limitation when it comes to storing large binary files, such as images, audio files, and video files, which can easily exceed the 16 MB document size limit.

To address this limitation, MongoDB provides GridFS, a specification that allows you to store and retrieve large files efficiently.

### 14.2 Using GridFS for file Management

GridFS stores large files as smaller, fixed-size chunks in MongoDB collections, making it possible to store and retrieve files that are much larger than 16 MB. GridFS also provides metadata storage,



which can include information like the file's name, content type, and additional attributes.

Here's how to work with GridFS in MongoDB:

## **1. Installing GridFS Drivers**

To work with GridFS, you'll need to install the MongoDB driver for your chosen programming language, as most drivers include GridFS functionality.

## **2. Uploading Files**

To upload a file using GridFS, you'll need to create a connection to your MongoDB database and specify the GridFS bucket where the file will be stored. Then, you can use the provided methods to upload the file.

**Example (Node.js with the “mongodb” driver):**

- **javascript**

```
const { MongoClient } = require('mongodb');
```

```
const fs = require('fs');
```

```
const uri = 'mongodb://localhost:27017';
```

```
const client = new MongoClient(uri);
```

```
async function uploadFile() {
```

```
  try {
```

```
    await client.connect();
```

```
    const database = client.db('mydatabase');
```

```
    const bucket = new database.GridFSBucket();
```

```
    const fileStream = fs.createReadStream('largefile.txt');
```

```
    const uploadStream = bucket.openUploadStream('largefile.txt');
```

```
    fileStream.pipe(uploadStream);
```

```
    console.log('File uploaded successfully');
```

```
  } finally {
```

```
    await client.close();
```

```
  }
```

```
}
```

```
uploadFile();
```

**In this example,** we use Node.js with the “mongodb” driver to upload a file named 'largefile.txt' to the GridFS bucket.

### 3. Downloading Files

To download a file from GridFS, you'll need to create a connection to your MongoDB database, specify the GridFS bucket, and use the provided methods to retrieve the file.

Example (Node.js with the `mongodb` driver):

#### - javascript

```
const { MongoClient } = require('mongodb');
```

```
const fs = require('fs');
```

```
const uri = 'mongodb://localhost:27017';
```

```
const client = new MongoClient(uri);
```

```
async function downloadFile() {
```

```
  try {
```

```
    await client.connect();
```

```
    const database = client.db('mydatabase');
```

```
    const bucket = new database.GridFSBucket();
```

```
    const downloadStream = bucket.openDownloadStreamByName('largefile.txt');
```

```

const fileStream =
fs.createWriteStream('downloaded_largefile.txt');

downloadStream.pipe(fileStream);

console.log('File downloaded successfully');

} finally {

  await client.close();

}

}

downloadFile();

```

**In this example,** we use Node.js with the `mongodb` driver to download a file named 'largefile.txt' from the GridFS bucket and save it as 'downloaded\_largefile.txt'.

#### 4. Deleting Files

Deleting a file from GridFS involves specifying the file's unique identifier (usually the ObjectId) and removing it from the GridFS bucket.

**Example (Node.js with the “mongodb” driver):**

- **javascript**

```
const { MongoClient } = require('mongodb');
```

```
const uri = 'mongodb://localhost:27017';
```

```
const client = new MongoClient(uri);
```

```
async function deleteFile(fileId) {
```

```
  try {
```

```
    await client.connect();
```

```
    const database = client.db('mydatabase');
```

```
    const bucket = new database.GridFSBucket();-
```

```
    await bucket.delete(fileId);
```

```
    console.log('File deleted successfully');
```

```
  } finally {
```

```
    await client.close();
```

```
  }
```

```
}
```

```
deleteFile('5f8e3d2151f0b9e14c7d9e35');
```

**In this example,** we use Node.js with the `mongodb` driver to delete a file from the GridFS bucket based on its unique identifiers.